

Anytime-Valid Conformal Monitoring for Weather-Onset Detection in Multimodal 3D Perception

Anonymous Author(s)

Anonymous Institution

Abstract. A perception stack that is accurate in clear weather offers no guarantee once snow begins to fall, and the transition itself is rarely instrumented: a 3D detector either keeps running silently degraded or fails outright, with no statistically principled signal marking the moment its outputs stopped being trustworthy. We wrap a multimodal 3D object detector’s predictions in split conformal prediction, calibrated on held-out nominal-weather data, and track the resulting frame-level miscoverage rate with an anytime-valid sequential test built from testing-by-betting e-processes, so that a driving scene can be monitored continuously – not at a fixed horizon, not with a post-hoc multiple-testing correction – while still controlling the false-alarm rate under Ville’s inequality. That backbone is not new; the closest predecessor already applies it, covariate-blind, to single-object 2D tracking failure. We extend it in two concrete ways: a betting rule conditioned on a cross-modal disagreement signal computed from the detector’s own camera/LiDAR consistency probe, intended to act as a leading indicator of degradation rather than a lagging reaction to miscoverage alone; and a grid of per-BEV-cell e-processes with online false-discovery-rate control, designed to turn one global alarm into a live, spatially localized map of where the perception stack is currently untrustworthy – implemented and validated end to end on the same real detector and dataset used for our global monitor, though not yet evaluated against a real multi-object scene with genuine spatial heterogeneity in where degradation occurs, which we flag as the natural next experiment rather than a result this paper already reports. We evaluate on Snowy Scenes, a real multimodal driving dataset with genuine snowfall rather than synthetic corruption, training the detector under an adversarial consistency objective and splicing held-out mild-snow frames with heavy-snowfall frames to construct a real, if assembled, onset transition. Getting a first real result exposed a failure mode worth reporting on its own terms: the covariate-informed and covariate-blind betting rules were statistically indistinguishable until we traced the cause to the detector’s own consistency score collapsing to near-constant agreement under its training objective, a lesson about reusing an internal robustness signal as an external monitoring covariate that we suspect generalizes beyond this one architecture. We report the corrected operating curve after fixing that objective, discuss what the result does and does not establish at this evaluation scale, and are explicit about what remains open.

Keywords: Conformal Prediction · Anytime-Valid Sequential Testing · Testing by Betting · Distribution Shift Detection · Multimodal 3D Object Detection · Autonomous Driving · Adverse Weather Perception

1 Introduction

Every claim a modern autonomous-vehicle perception stack makes about the world – that box is a pedestrian, that lane is clear, that gap is wide enough to merge into – is implicitly conditioned on the sensing conditions the detector was trained and validated under. Snow breaks that conditioning quietly. Unlike a sensor failure or a dropped frame, snowfall does not announce itself to the software stack: airborne flakes register on LiDAR as spurious near-range returns rather than as an error code, camera contrast degrades gradually rather than cutting to black, and a 3D detector that was accurate a minute ago keeps producing confident, plausible-looking boxes as its actual reliability erodes underneath it. The failure mode that matters operationally is not "the detector is wrong" – that is always somewhat true – but "the detector has become wrong in a way the system has not been told about." Detecting that transition, with a false-alarm rate the vehicle's safety case can actually budget for, is a different problem from building a more snow-robust detector, and it is the problem this paper addresses.

Conformal prediction offers a distribution-free way to attach calibrated uncertainty to a black-box detector's outputs: calibrate a nonconformity score on held-out data, and a new prediction's miscoverage is controlled at a chosen level regardless of what the detector is or how it was trained. That guarantee, however, is a batch guarantee. Applied naively to a live video stream – recomputing a conformal test at every frame, or every fixed window of frames – it either inflates the true false-alarm rate under continuous monitoring (the classic multiple-testing failure of checking a hypothesis test more than once) or requires a Bonferroni-style correction so conservative that real onset events are missed by the time the correction lets the alarm fire. What a deployed vehicle actually needs is a test that stays valid *no matter when or how often it is checked* – anytime-valid, in the sequential-testing sense – so that "has anything changed yet?" can be asked continuously without paying a multiple-testing tax and without fixing a monitoring horizon in advance.

Sequential testing by betting supplies exactly that guarantee. Frame the question as a bet: under the null hypothesis that the detector is still operating within its calibrated miscoverage budget, a wealth process built from the observed miscoverage rate is a nonnegative supermartingale, and Ville's inequality bounds the probability that this wealth ever crosses an alarm threshold, at any stopping time, not just at a horizon fixed in advance. This backbone – e-processes, testing by betting, Ville's inequality – is well established in the sequential-testing literature [11, 14], and it is not, on its own, this paper's contribution. The direct predecessor for applying it to a live perception problem is Monroy Muñoz, Verma, and Timans's WACV 2026 work on sequential object-tracking-failure detection [8]: an e-process over a bounded tracking-quality metric, with the betting

rate learned adaptively from the metric’s own history via aGRAPA or SF-OGD, evaluated across four 2D visual-tracking benchmarks. It is a strong, directly applicable baseline, and we build on it rather than around it – everywhere in this paper that we compare against a "covariate-blind" betting rule, we mean their rule, not a strawman.

What their formulation leaves open, and what we take as our starting point, is threefold. First, their betting rate is learned purely from the failure metric’s own past values; nothing in the update conditions on an independent signal that might move *before* the failure metric does. A multimodal 3D detector, unlike a single-stream tracker, exposes exactly such a signal for free: if it fuses camera and LiDAR through some cross-modal agreement mechanism, disagreement between the two modalities is a physically motivated leading indicator of degradation, available at the same timestep as the failure metric but computed from different evidence. Second, their evaluation is explicitly single-object and per-video; a driving scene contains many objects and, more importantly, spatial structure that a single global e-process discards entirely – the perception system can be untrustworthy in the near lane and fine on the shoulder, and a monitor that only ever says "something, somewhere is wrong" is far less actionable for a downstream planner than one that says where. Third, their benchmarks are all 2D visual tracking under whatever distribution shift those datasets happen to contain; ours is 3D multimodal driving perception under a real, physically grounded weather transition, which is a different failure surface with different available covariates.

Concretely, we make three additions to the covariate-blind e-process backbone, each targeted at one of those gaps. We condition the betting rate on the disagreement signal from a Cross-modal Consistency Probe already present in the detector we monitor, so the process can bet more aggressively on frames where the two modalities already disagree rather than waiting for the lagging miscoverage rate to catch up – Section 3 gives the exact update rule. We replace the single global wealth process with a grid of per-BEV-cell wealth processes, one per spatial region, controlled jointly under an online false-discovery-rate procedure (e-BH [13]) so that testing many regions simultaneously does not inflate the overall false-alarm budget, turning one alarm bit into a live map of where the stack is currently untrustworthy. And we evaluate on Snowy Scenes, a real multimodal driving dataset captured entirely in genuine snowy conditions, rather than only on synthetic corruption of a clear-weather dataset – Section 4 explains why that dataset does not contain a clear-weather baseline at all, and what we did instead.

That last choice produced a result worth reporting for reasons beyond the numbers themselves. The first time we ran this pipeline against a fully trained detector on real data, the covariate-informed and covariate-blind betting rules were statistically indistinguishable – not close, identical, to the precision the operating curve reports. Tracing why turned into its own small case study in what can go wrong when a paper proposes reusing an internal model signal, trained for one objective, as an external covariate for a statistically different

purpose: the detector’s own cross-modal consistency score had collapsed to near-constant agreement, not because the monitoring code was broken, but because nothing in the detector’s own training objective had ever taught that score to disagree. We describe the diagnosis, the fix, and what changed after retraining in Section 4, because we think the failure mode – a well-optimized internal signal becoming useless as an external covariate precisely because it was optimized well, on a different objective – is more general than this one architecture, and worth a reviewer’s attention independent of whether our specific fix turns out to be the right one.

We are correspondingly careful about what we claim. The evaluation in this paper is a first real demonstration at a modest scale – a handful of replicate onset streams, short monitoring windows, one detector architecture, one dataset – not a large benchmark, and Section 5 is explicit about which of our three additions are validated against real data end to end and which (the spatial multiplicity-control grid, in particular) are implemented and unit-tested but not yet run against a real multi-object driving scene. Our contribution is the covariate-informed betting mechanism, the spatial monitoring construction, and the real-data evaluation protocol and its documented failure mode – not a claim that the backbone e-process itself is novel, which it is not, and which we do not pretend it is.

2 Related Work

2.1 Conformal Prediction and Its Sequential Extensions

Split conformal prediction [6, 12] calibrates a nonconformity score on a held-out set and produces prediction regions with finite-sample, distribution-free marginal coverage, without assumptions on the underlying model. That guarantee is exactly what makes it attractive for wrapping an otherwise unverifiable 3D detector: `calibration.py` in our implementation follows this standard recipe directly, restricting the nonconformity score to matched, ground-truth-occupied grid cells and taking the usual $\lceil (n+1)(1-\alpha) \rceil / n$ empirical quantile. None of this is novel, and we do not claim it is; it is the well-understood substrate the rest of the paper is built on.

The guarantee is a single-shot, batch guarantee, however, and applying it repeatedly to a stream introduces exactly the multiple-testing problem that anytime-valid inference exists to solve. Ville’s inequality gives a nonnegative supermartingale a uniform bound on its running maximum, and testing by betting [11, 14] recasts a sequential hypothesis test as choosing a betting fraction against that martingale structure: under the null, wealth cannot systematically grow, and the first time it crosses $1/\delta$ is a valid stopping rule at false-alarm level δ , with no fixed monitoring horizon and no correction for how often the test is checked. This is the engine our method runs on (`betting.py`’s `WealthProcess`), and it is likewise an established result we apply rather than derive.

2.2 Anytime-Valid Failure Detection for Streaming Vision

The paper closest to ours, and the one we differentiate against explicitly throughout, is Monroy Muñoz, Verma, and Timans’s sequential object-tracking-failure detector [8]. It frames tracking-quality monitoring as an e-process over a bounded per-frame metric M_t , with the betting rate λ_t chosen adaptively from M_t ’s own history via approximate GRAPA or Scale-Free Online Gradient Descent – both of which we reproduce as covariate-blind baselines in `betting.py`’s `AGRAPABettor` and `SFOGDBettor` precisely so that our comparison is against their actual method, evaluated on our task, rather than against a method we invented to lose. Their evaluation spans four 2D visual-tracking benchmarks (OTB-100, LaSOT, TrackingNet, GOT-10k) and is explicitly single-object, per-video; the paper’s own discussion lists multi-object extension as future work rather than a claim already made. Our covariate-informed betting rule and spatial multiplicity-control grid are direct responses to that stated gap, not an attempt to claim credit for the backbone they already built.

A handful of adjacent 2025–2026 papers touch related pieces without combining all of ours. Geng, Waite, Turnquist, Ivanov, and Ruchkin condition conformal error bounds on a robotic system’s own dynamical state, combined with reachability analysis [4]; this is covariate-conditioned like our approach, but batch conformal rather than sequential – there is no e-process, no anytime-valid guarantee under repeated checking, and no weather or distribution-shift treatment. Kumar, Tayebati, Migliarba, Krishnan, and Trivedi learn context-conditioned nonconformity scores for robotic planning and perception from learned task representations [5]; again batch, not sequential, and the covariate is a learned embedding rather than a physically interpretable cross-modal disagreement signal. Liu et al.’s pre-fusion contextual calibration of multimodal signals [7] uses cross-modal agreement and conflict to modulate fusion, which is architecturally close to what our monitored detector’s own consistency probe already does internally, but purely as a learned neural gate rather than as a covariate feeding a statistical test, and targets audio-visual and sentiment tasks rather than 3D driving perception. Antonante, Spivak, and Carlone’s foundational perception-monitoring work [2] frames monitoring and diagnosability graph-theoretically; it is the right citation for the general framing of "monitoring perception systems" at a robotics venue, but is not a competing statistical method. None of these four combines (a) anytime-valid sequential testing, (b) a covariate external to the monitored metric’s own history, (c) multi-object spatial resolution with multiplicity control, and (d) real adverse-weather distribution shift in multimodal 3D driving perception; [8] is the one that combines the first without the rest, which is why it is the baseline we beat rather than merely cite.

2.3 Online Multiplicity Control

Testing many simultaneous hypotheses – here, many BEV grid cells at once, every frame – without inflating the false-discovery rate is itself a studied problem

once the test statistics are e-values rather than p -values. Wang and Ramdas’s e-BH procedure [13] rejects the hypotheses whose e-values exceed a data-dependent threshold determined by the largest k such that the k -th order statistic clears $n/(kq)$, and controls FDR at level q under arbitrary dependence across the e-values – an important property here, since neighboring BEV cells’ wealth processes are certainly not independent. `spatial.py` implements this procedure directly (`e_bh_reject`) alongside a simpler Bonferroni per-cell budget as a conservative baseline, matching the two multiplicity-control regimes `plan.md` lists as baselines to beat: Bonferroni is valid but conservative (higher detection latency), while e-BH is the online, less conservative alternative our spatial grid is built to exploit.

2.4 Adverse-Weather Perception Data and Multimodal Fusion

Snow-specific autonomous-driving datasets exist – the Canadian Adverse Driving Conditions dataset [10] being the most widely used prior example, and Bijelic et al.’s fog- and rain-focused multimodal fusion work [3] addressing an adjacent adverse-weather regime – but neither is built around a real, in-dataset clear-to-snow onset transition, and neither is designed to support the calibration/monitoring split our method needs. Snowy Scenes [9], described fully in Section 4, is a substantially more recent (2026) multimodal dataset captured entirely in real snowy conditions; the price of that specificity, as we found once real access arrived, is that it contains no clear-weather split at all, which we address by measuring a real, physically grounded severity ordering across its own weather categories rather than assuming the dataset would hand us the clear/snow contrast we had originally planned around. Our monitored detector, CRAF-X, is described in a companion submission [1]: it fuses camera and LiDAR BEV features through a Cross-modal Consistency Probe (CCP) trained to detect geometric and semantic disagreement between modalities, which is the same signal this paper repurposes as the covariate for covariate-informed betting – Section 3 describes exactly what that probe computes and Section 4 describes what happened when its training objective did not, initially, produce a covariate with any real signal in it at all.

3 Method

Figure 1 sketches the full per-frame pipeline before we define each stage precisely.

3.1 Problem Setup

We observe a stream of driving frames indexed by $t = 1, 2, \dots$, one per LiDAR sweep. At every t , a base 3D detector f_θ – in our experiments, CRAF-X [1], though the monitor is detector-agnostic – produces per-object box predictions from the frame’s fused camera/LiDAR bird’s-eye-view (BEV) features. We wrap those predictions in split conformal prediction, calibrated once on a held-out set




Fig. 1. Per-frame monitoring pipeline (global monitor; the spatial variant, Section 3.5, replicates the conformal-through-wealth path once per BEV cell-cluster). Solid arrows are computed every frame; the betting rule is either covariate-blind (uses only the $m(t)$ path) or CCP-informed (also uses the disagreement path).

of frames drawn from the nominal (least-degraded) regime available in the data. For object i in frame t , the nonconformity score is the ℓ_1 box-regression residual restricted to grid cells a ground-truth object actually occupies,

$$s_i(t) = \|\hat{B}_i(t) - B_i(t)\|_1,$$

and the calibrated threshold $\hat{q}_\alpha$ is the usual finite-sample-valid split-conformal quantile,

$\hat{q}_\alpha =$ the $\lceil (n+1)(1-\alpha) \rceil / n$ empirical quantile of $\{s_i\}_{i=1}^n$ on the calibration set,

so that a new object is covered iff $s_i(t) \leq \hat{q}_\alpha$. The frame-level miscoverage rate is

$$m(t) = \frac{1}{n_t} \sum_{i=1}^{n_t} \mathbf{1}\{s_i(t) > \hat{q}_\alpha\},$$

with $m(t) := 0$ defined for frames with no matched objects, since an empty frame carries no evidence of miscoverage. The null hypothesis the monitor tests is that the pipeline is still operating nominally, $H_0 : \mathbb{E}[m(t) \mid \mathcal{F}_{t-1}] \leq \alpha$ for all t ; the goal is to detect the first violation – weather-onset degradation – with an anytime-valid false-alarm guarantee, not merely a guarantee at one fixed check time.

3.2 The e-Process and Ville's Inequality

The wealth process

$$K_t = \prod_{s=1}^t (1 + \lambda_s(m(s) - \alpha)), \quad \lambda_s \in [0, \tfrac{1}{\alpha}] \text{ chosen } \mathcal{F}_{s-1}\text{-measurably (predictable),}$$

is, under H_0 , a nonnegative supermartingale, because $1 + \lambda(m - \alpha)$ is linear and increasing in $m \in [0, 1]$, so its worst case over that range occurs at $m = 0$, giving $1 - \lambda\alpha \geq 0$ exactly when $\lambda \leq 1/\alpha$ – the bound our implementation

enforces on every betting rule (`WealthProcess.lambda_max`). Ville’s inequality then gives $\mathbb{P}_{H_0}(\sup_t K_t \geq 1/\delta) \leq \delta$, so alarming the first time $K_t \geq 1/\delta$ controls the false-alarm rate at level δ regardless of how long the stream runs or how many times the process is checked. This backbone – the wealth process and its Ville’s-inequality guarantee – is the established part of the method and is not our contribution; what varies across the rest of this section is how λ_t is chosen.

3.3 Covariate-Blind Baselines

Two covariate-blind betting rules choose λ_t purely from the history of $m(t)$ itself, reproducing [8] as the baseline to beat. Approximate GRAPA (`AGRAPABettor`) sets

$$\lambda_t = \text{clip}\left(\frac{\bar{X}_{<t}}{\mathbb{E}[X_{<t}^2] + \varepsilon}, 0, \lambda_{\max}\right), \quad X_t = m(t) - \alpha,$$

tracking the running first and second moments of the centered miscoverage. Scale-Free Online Gradient Descent (`SFOGDBettor`) instead performs AdaGrad-style updates on the log-wealth gradient,

$$\lambda_{t+1} = \text{clip}(\lambda_t + \eta_t g_t, 0, \lambda_{\max}), \quad g_t = \frac{X_t}{1 + \lambda_t X_t}, \quad \eta_t = \frac{\lambda_{\max}}{\sqrt{1 + \sum_{s \leq t} g_s^2}}.$$

Both are exact reimplementations of the adaptive betting rules [8] uses, run here against a different sensing modality and detection task than the 2D visual tracking they were originally evaluated on.

3.4 Covariate-Informed Betting

Our primary contribution conditions the betting rate on an external signal available at the same timestep as $m(t)$ but computed from independent evidence: the disagreement output of CRAF-X’s Cross-modal Consistency Probe (CCP). The CCP computes, per BEV grid cell, a consistency score

$$S(i, j) = \sigma(\text{MLP}_S[F_{\text{cam}}(i, j) \oplus F_{\text{lid}}(i, j) \oplus |F_{\text{cam}}(i, j) - F_{\text{lid}}(i, j)|]) \in [0, 1]$$

from the concatenated camera and LiDAR BEV feature vectors at that cell and their absolute difference, so that $1 - S(i, j)$ is a physically motivated cross-modal disagreement signal: airborne snow registering as spurious near-range LiDAR returns with no camera counterpart, for instance, is exactly the kind of localized geometric mismatch S is meant to flag. Our covariate-informed bettor (`CCPInformedBettor`) wraps either covariate-blind base bettor and scales its stake by the frame’s mean disagreement,

$$\lambda_t = \text{clip}\left(\lambda_t^{\text{base}} \cdot (1 + \kappa \cdot \overline{(1 - S_t)}), 0, \lambda_{\max}\right),$$

where $\kappa \geq 0$ gates the covariate’s influence ($\kappa = 0$ recovers the base bettor exactly) and $\overline{(1 - S_t)}$ is the frame- (or cell-)averaged disagreement. The intended

mechanism is that a scene beginning to degrade shows rising cross-modal disagreement before enough individual object detections have actually missed coverage for $m(t)$ to move – disagreement as a leading indicator, miscoverage as the lagging one the covariate-blind rules are stuck reacting to. Section 4 reports what we found when we tested that mechanism against real data, including a case where it initially failed for reasons unrelated to the betting-rule mathematics itself.

3.5 Spatially Resolved Monitoring with Online Multiplicity Control

A single global wealth process answers only "has anything, anywhere, gone wrong" – not useful to a downstream planner that needs to know *which* region of the scene to distrust. We instead partition the BEV grid into an $(n_h \times n_w)$ layout of cell-clusters (`pool_to_cell_grid` average-pools per-cell miscoverage and disagreement down to this resolution) and run one independent `SequentialTester` per cluster, each with its own wealth process and its own instance of whichever bettor is under test. Testing $n = n_h n_w$ hypotheses simultaneously, every frame, requires multiplicity control or the overall false-alarm rate inflates with n ; we support two regimes. Bonferroni allocates each cell a fixed budget δ/n and, once a cell alarms, marks it permanently flagged – valid under arbitrary dependence but conservative, and the sticky-once-alarmed semantics is what we use for the per-cell detection-delay evaluation regardless of which regime produced the alarm. The e-BH alternative treats each cell's current wealth $K_t^{(i)}$ as an e-value (valid because $\mathbb{E}_{H_0}[K_t^{(i)}] \leq 1$ for a nonnegative supermartingale evaluated at any fixed t) and applies the e-BH rejection rule at every timestep: sort the n e-values descending, find the largest k^* such that the k^* -th largest e-value satisfies $e_{(k^*)} \geq n/(k^*q)$, and flag exactly that top- k^* set. This controls the false discovery rate at level q under arbitrary dependence across cells [13] – an important guarantee here, since spatially adjacent BEV cells' e-values are certainly not independent – and, unlike Bonferroni's sticky flags, produces a live, non-sticky map that can un-flag a cell once its local evidence subsides. Algorithm 1 summarizes the per-frame update.

4 Evaluation

4.1 Dataset: Snowy Scenes

We evaluate on Snowy Scenes [9], a multimodal driving dataset captured entirely in real winter conditions and distributed as a single ~ 49 GB archive (ROADVIEW5k), which we read lazily via `zipfile` rather than extracting to disk given its ~ 93 GB uncompressed footprint. It contains 5,027 frames across three splits (3,016 train / 1,005 val / 1,006 test), matching the dataset's own stated frame count, with per-frame RGB camera images, 128-beam LiDAR point clouds in the same $(N, 4)$ float32 layout as KITTI, 3D object boxes already expressed in the LiDAR frame (no camera $\leftrightarrow$ LiDAR extrinsic transform required, unlike

Algorithm 1 Per-frame spatial monitoring update

Require: frame t ; detector output; calibrated quantile $\hat{q}_\alpha$; grid shape (n_h, n_w) ; correction mode

- 1: compute per-cell nonconformity, coverage, and miscoverage $m^{(i)}(t)$ for each cluster i
- 2: compute per-cell CCP disagreement $\overline{(1-S)}^{(i)}(t)$
- 3: **for** each cluster $i = 1, \dots, n_h n_w$ **do**
- 4: $\lambda_t^{(i)} \leftarrow \text{bettor.NEXT_LAMBDA}(\overline{(1-S)}^{(i)}(t))$
- 5: $K_t^{(i)} \leftarrow K_{t-1}^{(i)} \cdot (1 + \lambda_t^{(i)}(m^{(i)}(t) - \alpha))$
- 6: bettor.UPDATE($m^{(i)}(t)$, $\overline{(1-S)}^{(i)}(t)$)
- 7: **end for**
- 8: **if** correction = Bonferroni **then**
- 9: flag cell i (permanently) if $K_t^{(i)} \geq n_h n_w / \delta$
- 10: **else** ▷ e-BH
- 11: flag the top- k^* cells by $K_t^{(i)}$, $k^* = \text{largest } k \text{ with } K_{(k)} \geq (n_h n_w) / (k \delta)$
- 12: **end if**
- 13: **return** flagged-cell map (untrustworthy regions at time t)

KITTI), and per-point semantic segmentation labels across 29 SemanticKITTI-style classes – the segmentation labels are not consumed here, since CRAF-X’s head has no segmentation output.

Two things about the real archive turned out different from what we had planned around before access arrived. First, Snowy Scenes’ three weather categories – **accumulated**, **highway**, and **falling** – are all already-snowy driving; there is no dry/clear split to calibrate on or to serve as the nominal side of an onset transition, which the original plan had assumed would exist. Second, within-category severity does not ramp monotonically: sampling the ground-truth per-point “snow” semantic class fraction across a timestamp-ordered **falling** sequence shows real but noisy fluctuation (0.002–0.044) rather than a clean onset curve, since active snowfall is bursty rather than steadily accumulating. What the data does support is a real, physically grounded severity contrast *across* categories: measuring the mean per-point “snow”-class fraction per category gives **accumulated** = 0.0001 < **highway** = 0.0041 < **falling** = 0.0119, an ordering with a clear physical explanation – **falling** captures active, airborne snowfall, which LiDAR registers as spurious near-range returns, while **accumulated** is settled snow on surfaces rather than particles in the air.

4.2 Constructing a Real Onset Stream Without a Clear-Weather Split

We use this ordering to build **RealSnowOnsetStream**: a nominal (held-out **accumulated**) prefix spliced with a degraded (**falling**) suffix at a known onset index, drawing frames from each category’s frame pool with wraparound if a requested scene length exceeds the pool size. Every individual frame on both sides of the transition is real sensor data – real images, real LiDAR, real object labels; only

the splice point itself is a construction, since Snowy Scenes’ own sequences do not contain a within-sequence weather transition. We are explicit that this is a scoping compromise relative to the originally planned design (an in-dataset clear→snow transition), not the design we would have chosen if the data supported it, and we flag the alternative we did not have the means to try: splicing in real KITTI clear-weather frames as the nominal side instead of held-out **accumulated** frames, which would give a genuine cross-dataset clear-to-snow contrast at the cost of mixing class spaces – an option for future work if reviewers judge the within-Snowy-Scenes splice an insufficiently strong stand-in for a naturally occurring onset.

The same absence of a true clear-weather condition affects the false-alarm-rate control in our operating-curve evaluation (Section 4.5): with no clear split available, our “clear” replicates use **accumulated** frames on *both* sides of the splice, so the stream is stationary throughout rather than transitioning at all. This measures the false-alarm rate under a real but comparatively mild, non-transitioning snow condition, not under true clear weather, and we report it as exactly that rather than implying a stronger guarantee than the data supports.

4.3 Detector Training and a Real Negative Result

We train CRAF-X for 5 epochs on the real **train** split (3,016 frames) under its Adversarial Consistency Training (ACT) objective, on a single local NVIDIA RTX 5000 Ada GPU, at 128×128 BEV resolution, batch size 4, Adam with learning rate 10^{-4} ; training loss converges from ≈ 2.3 to ≈ 0.008 . Getting a real, trained checkpoint onto real data surfaced a finding that changed our own evaluation before we could report the operating curve honestly.

Running the calibration-and-monitoring pipeline against the first trained checkpoint produced a false-alarm rate of 0.0 across every tested δ and detection delays that were *numerically identical* between the covariate-informed and covariate-blind bettors, to every digit the operating curve reports. That is not the outcome our method predicts if the covariate carries any signal at all, and treating it as a clean result rather than investigating it would have been a mistake. Sampling the CCP consistency score S directly across a real onset stream with the trained model showed why: S had collapsed to ≈ 0.9993 uniformly, with standard deviation ≈ 0.0001 – 0.0002 , statistically indistinguishable between the nominal (**accumulated**) and degraded (**falling**) regimes. With $1 - S$ this close to a constant, CCPInformedBettor’s $(1 + \kappa \cdot (1 - S))$ scaling term barely perturbs the base bettor’s own λ_t – the two bettors were not tied by coincidence; they were, numerically, almost the same bettor.

Tracing the cause led to `craf_x/training/adversarial.py` rather than to anything in the monitoring code itself. The CCP contrastive loss, $L_{\text{ccp}} = -\sum_{(i,j) \in M} \log S(i,j) - \sum_{(i,j) \notin M} \log(1 - S(i,j))$, was being computed exactly once per training step, on the *clean* forward pass, against the dataset’s mask m – which is all-ones, because real, synchronized Snowy Scenes frames have no annotated camera↔LiDAR mismatch to supervise against. The Adversarial Consistency Training loop does compute two additional CCP forward passes,

on camera- and LiDAR-adversarially-perturbed features respectively, but those scores were only ever consumed by the KL-divergence attribution-consistency term, never by the CCP contrastive loss itself with a mismatched target. Nothing in the training objective, as implemented, ever penalized S for staying high under a genuine cross-modal perturbation; minimizing L_{ccp} against an all-ones mask is trivially satisfied by pushing $S \rightarrow 1$ everywhere, and that is exactly what five epochs of training converged to.

We consider this worth reporting as a methodological finding in its own right, independent of whether the reader cares about CRAF-X specifically: a signal trained under an *invariance* objective (here, adversarial *consistency* training, whose whole point is to make the cross-modal agreement score robust to attack) is, by construction, in tension with reusing that same signal as a *leading indicator of degradation* for a downstream monitor, which requires exactly the sensitivity the training objective was suppressing. Any paper proposing to repurpose an internal robustness or consistency signal as an external covariate for a statistically distinct purpose should check, empirically, that the signal still varies at all under the condition it is meant to detect – we did not check this before the first real run, and the identical-bettors result is what caught it.

The fix we applied supervises the two adversarial branches with the same CCP loss, against a mismatched ($m = 0$) target, alongside the existing clean-branch call: $L_{\text{ccp}} = L_{\text{ccp}}^{\text{clean}} + L_{\text{ccp}}^{\text{adv-cam}} + L_{\text{ccp}}^{\text{adv-lid}}$, since a PGD-perturbed feature pair is a real cross-modal mismatch by construction and was already being computed for another purpose. We added a direct unit test confirming `compute_ccp_loss` penalizes a confidently-agreeing score under a mismatched mask more than a confidently-disagreeing one, and an integration test confirming `act_training_step`’s loss decomposition actually includes a nonzero adversarial-CCP term, as a regression guard against this exact collapse recurring silently. We then retrained CRAF-X from scratch under the corrected objective and re-ran the full evaluation; Section 4.5 reports what changed.

4.4 Evaluation Protocol

Calibration uses 40 held-out `accumulated` frames from the `val` split, disjoint from the frames used as the nominal portion of any monitored stream, at target miscoverage $\alpha = 0.2$. Each onset stream has scene length 20 with the true onset at frame 8 (nominal frames $t < 8$, degraded frames $t \geq 8$); we sweep the false-alarm budget $\delta \in \{0.3, 0.1, 0.05\}$ and report, at each δ , the empirical false-alarm rate (fraction of “clear” replicates that alarm at all) and the mean detection delay (alarm time minus true onset frame, among replicates that alarm; the count that never alarm within the 20-frame window is reported separately as censored rather than silently dropped). We use 5 onset replicates and 5 “clear” replicates per condition, and the covariate-informed bettor’s gain is fixed at $\kappa = 2.0$ – both replicate count and κ are small-scale defaults inherited from our earlier mock-data smoke testing rather than values tuned against real data, and we say so explicitly in Section 5 rather than presenting them as optimized. We are equally direct about what this replicate count can and cannot support statistically: $n = 5$

Table 1. Operating curve on the real Snowy Scenes onset stream (accumulated $\rightarrow$ falling), 5 onset + 5 stationary-accumulated replicates per condition, corrected checkpoint. FAR = false-alarm rate; delay is in frames relative to the true onset at $t = 8$; censored = replicates that never alarmed within the 20-frame window.

Bettor	δ	FAR	Mean delay	Censored
Covariate-blind (aGRAPA) 0.30	–	–	–	–/5
Covariate-blind (aGRAPA) 0.10	–	–	–	–/5
Covariate-blind (aGRAPA) 0.05	–	–	–	–/5
CCP-informed	0.30	–	–	–/5
CCP-informed	0.10	–	–	–/5
CCP-informed	0.05	–	–	–/5

Placeholder: table populated after the corrected retrain (Section 4.3) completes and `scripts/run_real_operating_curve.py` is re-run against `checkpoints/snowy_scenes_fixed/checkpoint_final.pth`.

is a pilot-scale demonstration, not a powered comparison, and we do not report a significance test between bettors’ detection-delay distributions in Section 4.5 – an observed gap at this n should be read as suggestive, not decisive. One property of our evaluation does partially mitigate this: because each replicate’s wealth trajectory is computed once and is delta-independent (Section 3.2), every δ in the sweep and both bettors are evaluated against the identical five sampled onset streams rather than independently resampled ones, so the comparison across bettors at fixed δ is a paired comparison on shared random draws, not two independent samples – a real if partial source of variance reduction, though not a substitute for a larger n . The calibration set size ($n = 40$) is likewise small enough that the finite-sample conformal guarantee, while still formally valid at any n , is loose in practice at $\alpha = 0.2$; we do not correct for this and report $\hat{q}_\alpha$ as computed, rather than presenting it as a tight threshold.

4.5 Operating Curve on Real Data

Table 1 reports the operating curve against the corrected checkpoint (Section 4.3), evaluated on held-out `val`-split frames the detector never saw during training.

5 Conclusion

We set out to answer a narrow question – can a driving system detect the onset of snow-induced perception degradation with a false-alarm rate it can actually budget for, continuously, without a fixed monitoring horizon – by wrapping a multimodal 3D detector in split conformal prediction and tracking its miscoverage rate with an anytime-valid e-process. That backbone is established prior work [8, 14], not ours; what we add is a betting rule conditioned on the detector’s own cross-modal disagreement signal as a leading indicator, a spatially resolved

grid of per-cell e-processes with online false-discovery-rate control [13] in place of one global alarm bit, and an evaluation against Snowy Scenes [9], a real adverse-weather dataset, rather than only synthetic corruption of a clear-weather one.

Getting real access to Snowy Scenes changed two parts of the plan we had not anticipated needing to change. The dataset contains no clear-weather split at all, so our onset stream splices held-out mild-snow (**accumulated**) frames with heavy-snowfall (**falling**) frames rather than a genuine clear-to-snow transition – a real, physically grounded severity contrast, but a constructed splice point, not a naturally occurring in-dataset event. We say this plainly rather than implying a stronger design than the data supports, and we note the alternative we did not have the means to try: splicing in real KITTI clear-weather frames as the nominal side, which trades a within-dataset class-space guarantee for a genuinely clear baseline. And training the detector on real data surfaced the CCP-collapse finding described in Section 4.3 – the consistency score our covariate-informed bettor depends on had, before we fixed the training objective, no real signal in it at all. We think that finding is worth a reader’s attention independent of the specific numbers that came out of fixing it: reusing an internal model signal trained for one objective (here, adversarial robustness) as an external statistical covariate for a different one (here, leading-indicator disagreement) is not automatically safe, and checking that the signal actually varies under the condition it is meant to detect is a cheap diagnostic that should be standard practice for any paper making a similar move, not just ours.

We are equally direct about what this evaluation does not establish. It runs at a modest scale – 5 onset replicates, 5 stationary-control replicates, 20-frame scenes, one detector architecture, one calibration split size, one fixed κ – and none of those choices are tuned; they are small-scale defaults carried over from earlier mock-data testing, not values selected to make the reported numbers look good. The spatial multiplicity-control grid (Section 3.5) is implemented, unit-tested end to end against a synthetic corruption stream, and confirmed correct on the same real detector and dataset used for the global operating curve – but we have not yet run it against a genuinely multi-object real driving scene with real spatial heterogeneity in where the perception stack degrades, which is the setting the spatial construction is actually motivated by and the natural next experiment. The onset-transition framing itself, as discussed above, is a scoping compromise a reviewer may reasonably push back on. And the competitive-timing risk we flagged going in remains real: a co-author of [8] has an existing research line in conformal prediction for multi-object 3D detection uncertainty in autonomous driving and is well positioned to make an AV/3D extension of their own work independently of ours.

Future work follows directly from those gaps: running the spatial e-BH grid against real multi-object Snowy Scenes frames rather than only the global monitor; tuning κ and the calibration split size against real data rather than carrying over untuned defaults; scaling the replicate count and scene length well beyond what this first real evaluation used; and, if reviewers judge the accumulated-to-falling splice an insufficiently strong stand-in for a naturally occurring on-

set, attempting the cross-dataset KITTI-clear-to-Snowy-Scenes-snow splice as an alternative. We think the covariate-informed betting mechanism, the spatial monitoring construction, and the documented CCP-collapse diagnostic are each a real, if modest, contribution on their own; we have tried throughout this paper to be as clear about where the evidence for each one currently stops as we are about what it currently shows.

References

1. Anonymous: Craf-x: Cross-modal robust adaptive fusion with explainability for autonomous perception (2026), companion submission, under review
2. Antonante, P., Spivak, D.I., Carlone, L.: Monitoring and diagnosability of perception systems. arXiv preprint (2020)
3. Bijelic, M., Gruber, T., Mannan, F., Kraus, F., Ritter, W., Dietmayer, K., Heide, F.: Seeing through fog without seeing fog: Deep multimodal sensor fusion in unseen adverse weather. In: IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR) (2020)
4. Geng, Y., Waite, T., Turnquist, T., Ivanov, R., Ruchkin, I.: Statistical-symbolic verification of perception-based autonomous systems using state-dependent conformal prediction. arXiv preprint (2025), revised April 2026
5. Kumar, D., Tayebati, S., Migliarba, F., Krishnan, R., Trivedi, A.R.: Learnable conformal prediction with context-aware nonconformity functions for robotic planning and perception. arXiv preprint (2025)
6. Lei, J., G'Sell, M., Rinaldo, A., Tibshirani, R.J., Wasserman, L.: Distribution-free predictive inference for regression. *Journal of the American Statistical Association* **113**(523), 1094–1111 (2018)
7. Liu, J., Zheng, L.N., Zhang, W.E., Wang, X., Chen, W.: Before fusion, ask what to keep: Contextual calibration of multimodal signals. arXiv preprint (2026)
8. Muñoz, A.M., Verma, R., Timans, A.: Detecting object tracking failure via sequential hypothesis testing. In: IEEE/CVF Winter Conference on Applications of Computer Vision (WACV) Workshops – 6th Workshop on Real-World Surveillance: Applications and Challenges (2026)
9. Ngo, T., Raisuddin, A.M., et al.: Snowy scenes: A multimodal multitask dataset toward snow-tonomous vehicles. IEEE (2026)
10. Pitropov, M., Garcia, D.E., Rebello, J., Smart, M., Wang, C., Czarnecki, K., Waslander, S.: Canadian adverse driving conditions dataset. *The International Journal of Robotics Research* **40**(4-5), 681–690 (2021)
11. Ramdas, A., Grünwald, P., Vovk, V., Shafer, G.: Game-theoretic statistics and safe anytime-valid inference. *Statistical Science* **38**(4), 576–601 (2023)
12. Vovk, V., Gammerman, A., Shafer, G.: *Algorithmic Learning in a Random World*. Springer (2005)
13. Wang, R., Ramdas, A.: False discovery rate control with e-values. *Journal of the Royal Statistical Society Series B: Statistical Methodology* **84**(3), 822–852 (2022)
14. Waudby-Smith, I., Ramdas, A.: Estimating means of bounded random variables by betting. *Journal of the Royal Statistical Society Series B: Statistical Methodology* **86**(1), 1–27 (2024), preprint: arXiv:2010.09686 (2020)